


# Cosmological Simulations: Under the Hood


**Salman Habib**

Argonne National Laboratory  
Kavli Institute for Cosmological Physics


Standard Model



No dark energy



Warm dark matter



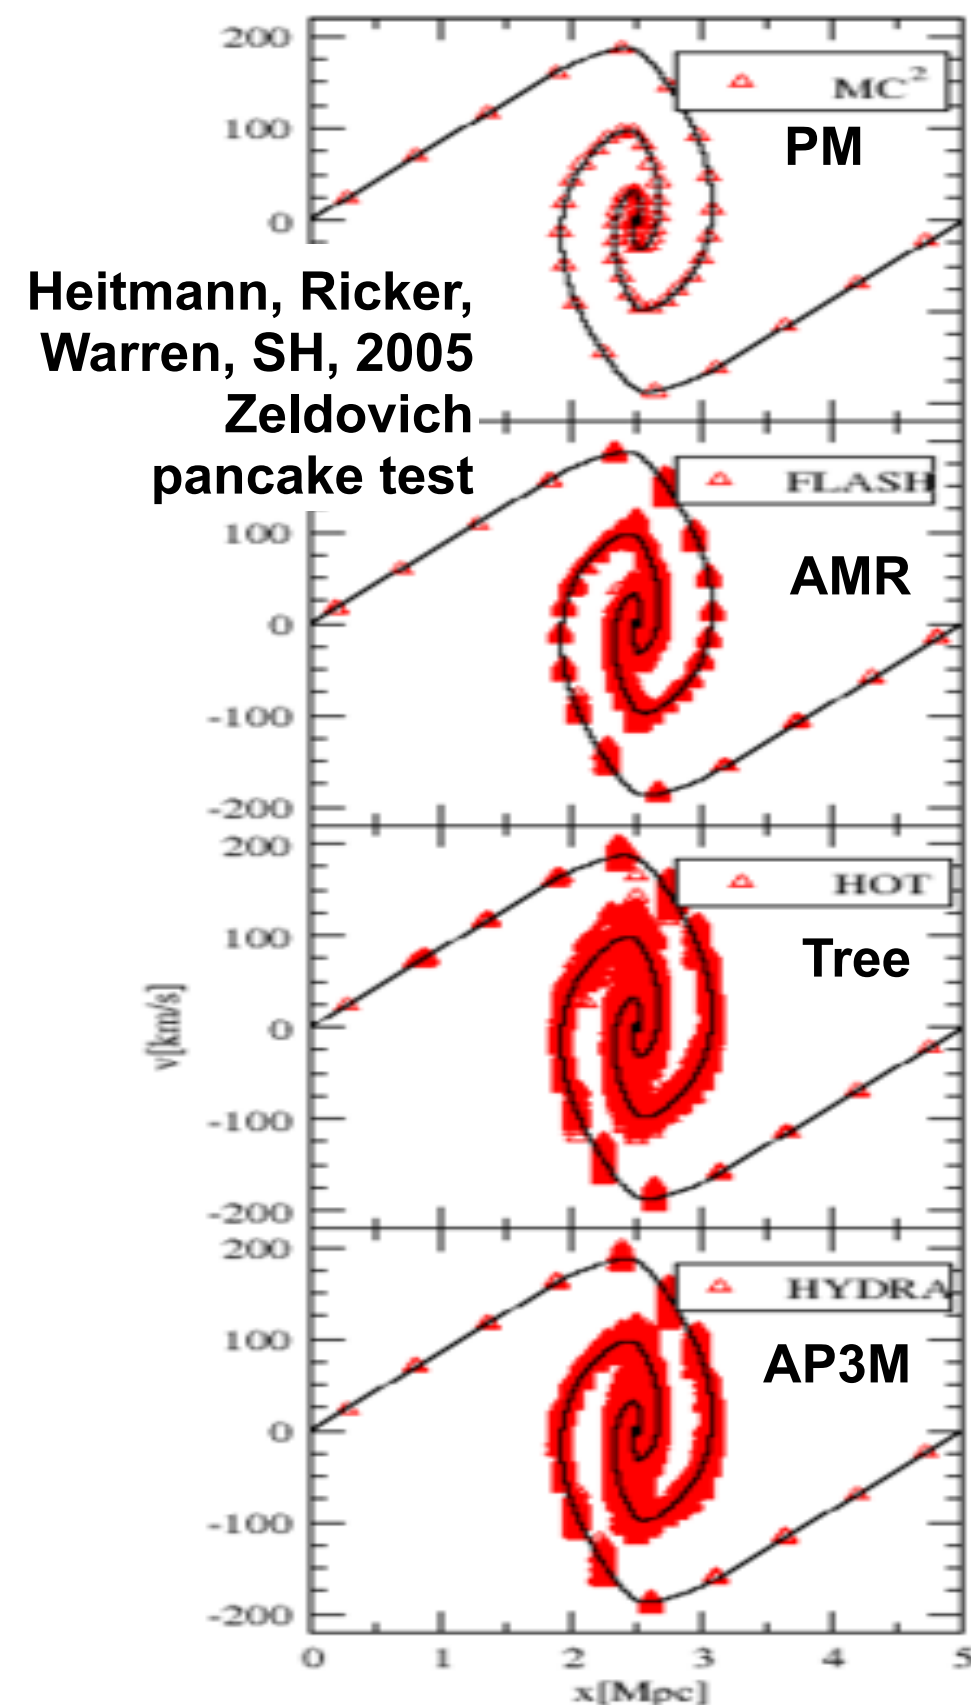
[http://www.nobelprize.org/nobel\\_prizes/physics/laureates/2016/advanced-physicsprize2016.pdf](http://www.nobelprize.org/nobel_prizes/physics/laureates/2016/advanced-physicsprize2016.pdf)

Haldane, Kosterlitz, Thouless (topological phase transitions)  
Home assignment — read the description of their work!

Mexican Numerical Simulations School  
Lecture 2, October 4, 2016

# Particle Approaches to the VPE: Overview

- **General remarks about particle methods**
  - Note that we are interested in a kinetic description because we have a collisionless system — the obvious fluid description is simply wrong (CDM is not an ideal gas!)
  - Later on we will like to combine the dark matter description with a separate Euler description for the baryonic fluid — how to best combine these two descriptions? (also, don't forget neutrinos!)
  - Because particle methods are intrinsically discrete, error analysis is subtle — many things can go wrong, and they often do!
  - Sometimes the best way to test results in complex simulation problems is to run multiple algorithms and compare results, so it is important to not focus on a single technique too much, but develop a suite of methods that have reasonable overlap in their domains of validity





# Particle Approaches: PM Method

- **Particle-In-Cell (PIC)/Particle-Mesh (PM)**
  - Reminder: Use tracer particles to model *collective* effects
  - Sequence of events: 1) generate ICs (particle positions and velocities), 2) generate density field on a grid, 3) solve the Poisson equations, generate gradient of the potential, 4) move the particles, 5) repeat
  - Note all information is particle information, the grid is a temporary construct to smooth the particle distribution and to compute the smooth force (well, more or less smooth)
  - Different particle deposition and force interpolation strategies (NGP, CIC, TSC); note: need symmetry in the deposition and interpolation schemes to have explicit momentum conservation
  - Different levels of smoothing allow different orders of Poisson solvers, typically accuracy and spatial resolution are in opposition (this may seem counter-intuitive, but we will see why it is so)
  - If at fixed particle number, one increases the grid size arbitrarily, one gets back an “N-body” problem, what is the correct choice of  $N_p$  vs.  $N_g$ ?

# Particle Approaches: P3M Method

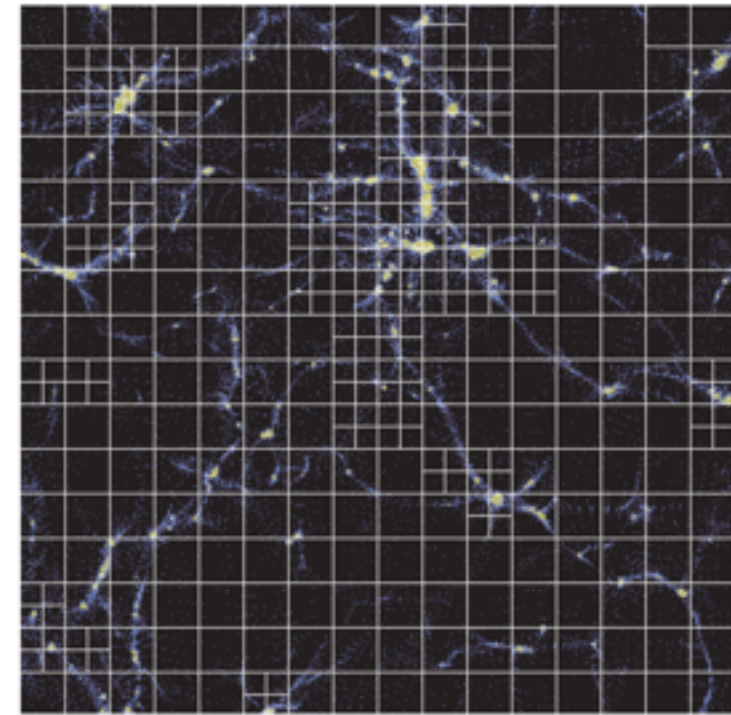
- **Particle-Particle Particle-Mesh (P3M)**

- Fundamental problem of the PM method is the memory cost of the grid
- If we can have a sufficiently large number of particles, increasing  $N_g$  to get enhanced resolution is potentially very expensive
- In P3M, one splits the force computation into two parts, a long-range force computed via PM (which also leaks into small scales) and a short-range force computed via direct particle-particle interactions
- Need to introduce a force smoothing scale for the particle-particle interaction to make sure we are still in the VPE limit (this is messy)
- To basic PM need to add another construct to reduce the particle-particle computational costs, the chaining mesh
- The particle-particle computations are expensive — for clustered problems, P3M can be potentially problematic
- Until recently, the general view was that P3M is not competitive for cosmological simulations, but GPUs have changed this picture


# Particle Approaches: AMR Techniques

## Adaptive Mesh Refinement (AMR)

- Fundamental problem of the PM method is the memory cost of the grid
- AMR attacks this by changing the size of the grid depending on the mass distribution
- Need to understand how different AMR levels interact
- Need to figure out a criterion for deciding the level of refinement (nontrivial, see bottom figure)
- More complex data structures needed
- Particle-mesh interaction complex (variable softening)
- Need a multi-scale Poisson solver, unlike PM or P3M.
- In high-resolution cosmological problems, because the resolution is needed “everywhere”, deep AMR requirements lead to high memory requirements
- Currently, deep AMR methods are mostly used for cosmological hydrodynamics simulations




**FLASH AMR hierarchy showing first two levels only, halo mass fn AMR issue below**



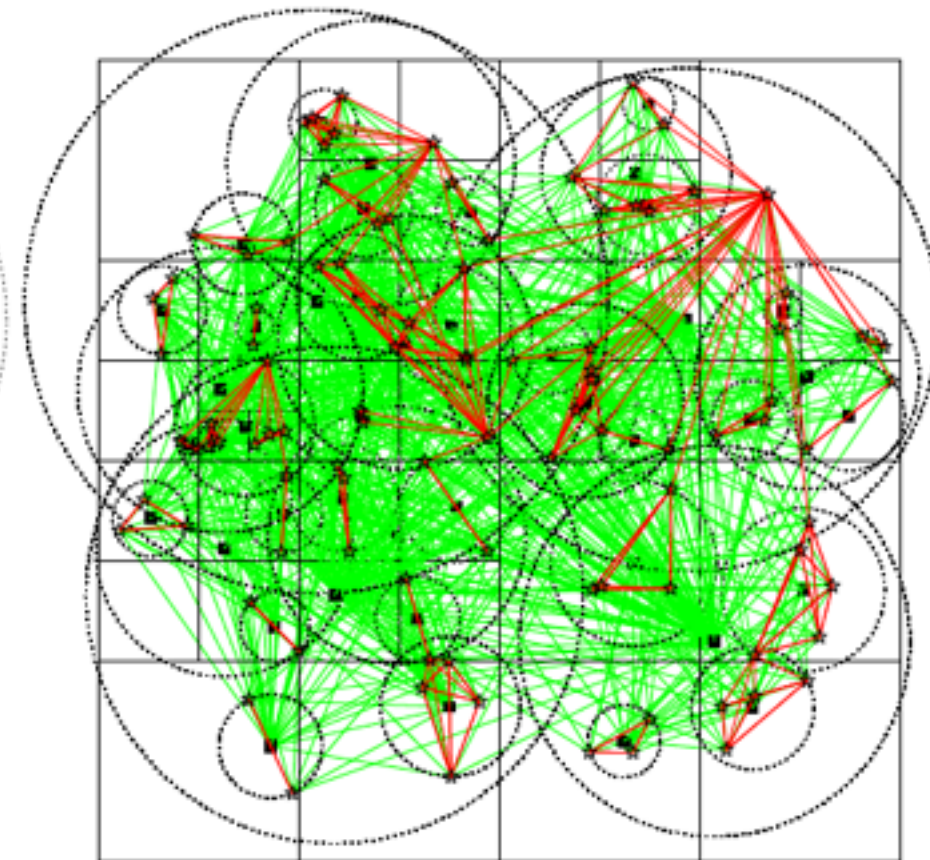
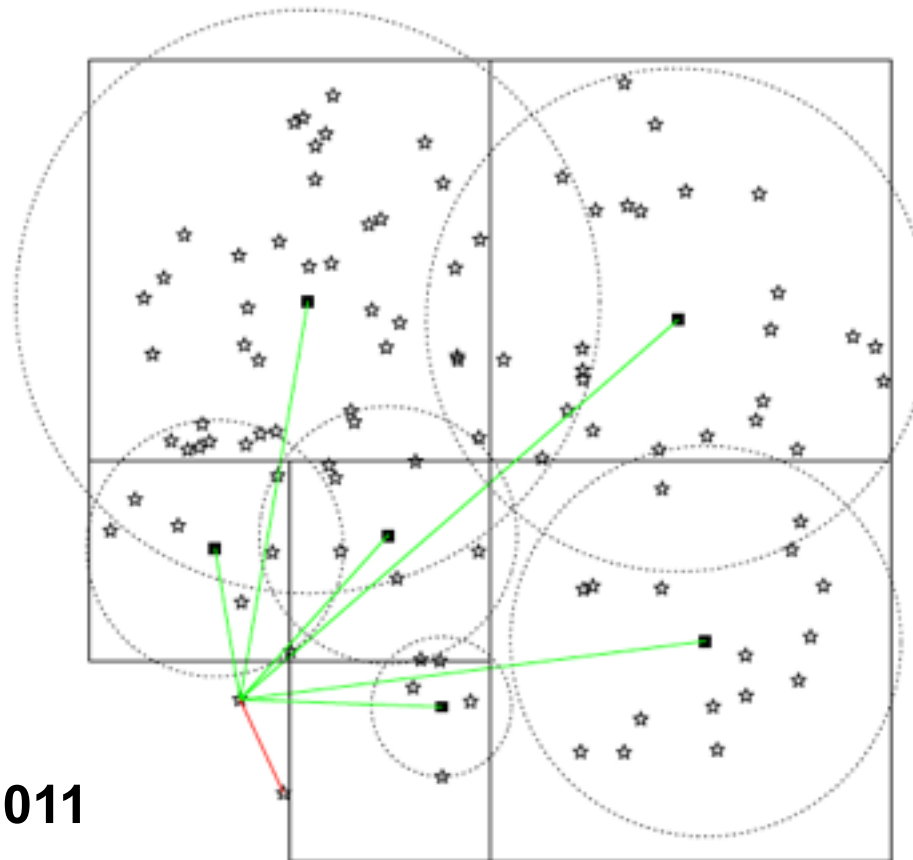


# Particle Approaches: Tree and FMM Techniques

- **Tree and Fast Multipole Algorithms**
  - Avoid grids altogether and exploit multipole expansions (particularly useful for clustered situations)
  - Need error control criteria (e.g., opening angle, expansion order) and appropriate data structures (RCB trees, oct-trees, space-filling curves, etc.)
  - Fast and efficient (FMM has a double expansion of the Green's function)
  - Not naturally periodic, need to add periodic BCs via Ewald sums




Dehnen and Read, 2011



# Particle Approaches: Hybrid Methods

- **Mostly PM + X (X=Tree or FMM)**
  - At large scales, PM methods are very convenient and fast, also good for evolving cosmological simulation at early times when clustering is small
  - Address weakness of PM codes at small scales via tree/FMM algorithms
  - Need to use force matching as in P3M but can be more relaxed because the efficiency of tree methods allows the matching point to be at a larger distance scale
  - HACCC uses PM + X (low-order FMM) to address multiple architectures




**Comparison of multiple algorithms on the same problem, halos are color-coded according to the number of particles (blue~100, green~1000, red~10,000)**

**SH et al. 2016**

# Case Study: PM

- **Particle deposition (density field)**
  - Use various weighting schemes to divide up the particle mass on nearby grid points (NGP — nearest grid point, CIC — volume-weighting to 8 grid points, TSC — quadratic on three nearest cells, 27 grid points)
  - NGP (field discontinuous), CIC (field continuous, gradient discontinuous), TSC (field and gradient both continuous)
  - NGP is rarely used (too noisy), CIC is most common
  - In Fourier space, the density weighting corresponds to sinc filters



**1-D representation of NGP, CIC, and TSC;  $h$  is the grid spacing**

$$\left[ \sin(\pi k \Delta / L) / (\pi k \Delta / L) \right]^n$$

**In Fourier space,  $n=1$  is NGP,  
 $n=2$  is CIC,  $n=3$  is TSC**



# Case Study: PM (FFT-Based Poisson Solver)

- **Poisson equation on a grid**

- 1-D: 
$$\frac{\phi(x + \Delta) + \phi(x - \Delta) - 2\phi(x)}{\Delta^2} = \frac{\partial^2 \phi}{\partial x^2} + O(\Delta^2)$$

- Use trigonometric collocation to show that the Influence function is

$$G_2(k) = \frac{\Delta^2}{2} \frac{1}{\cos(2\pi k \Delta / L) - 1}$$

- Note that this is “hotter” than the continuous Influence function, but it only appears in the Fourier integral multiplied by the CIC filtered density

- It is trivial to show that  $G_2(k)$  multiplied by the CIC filter is

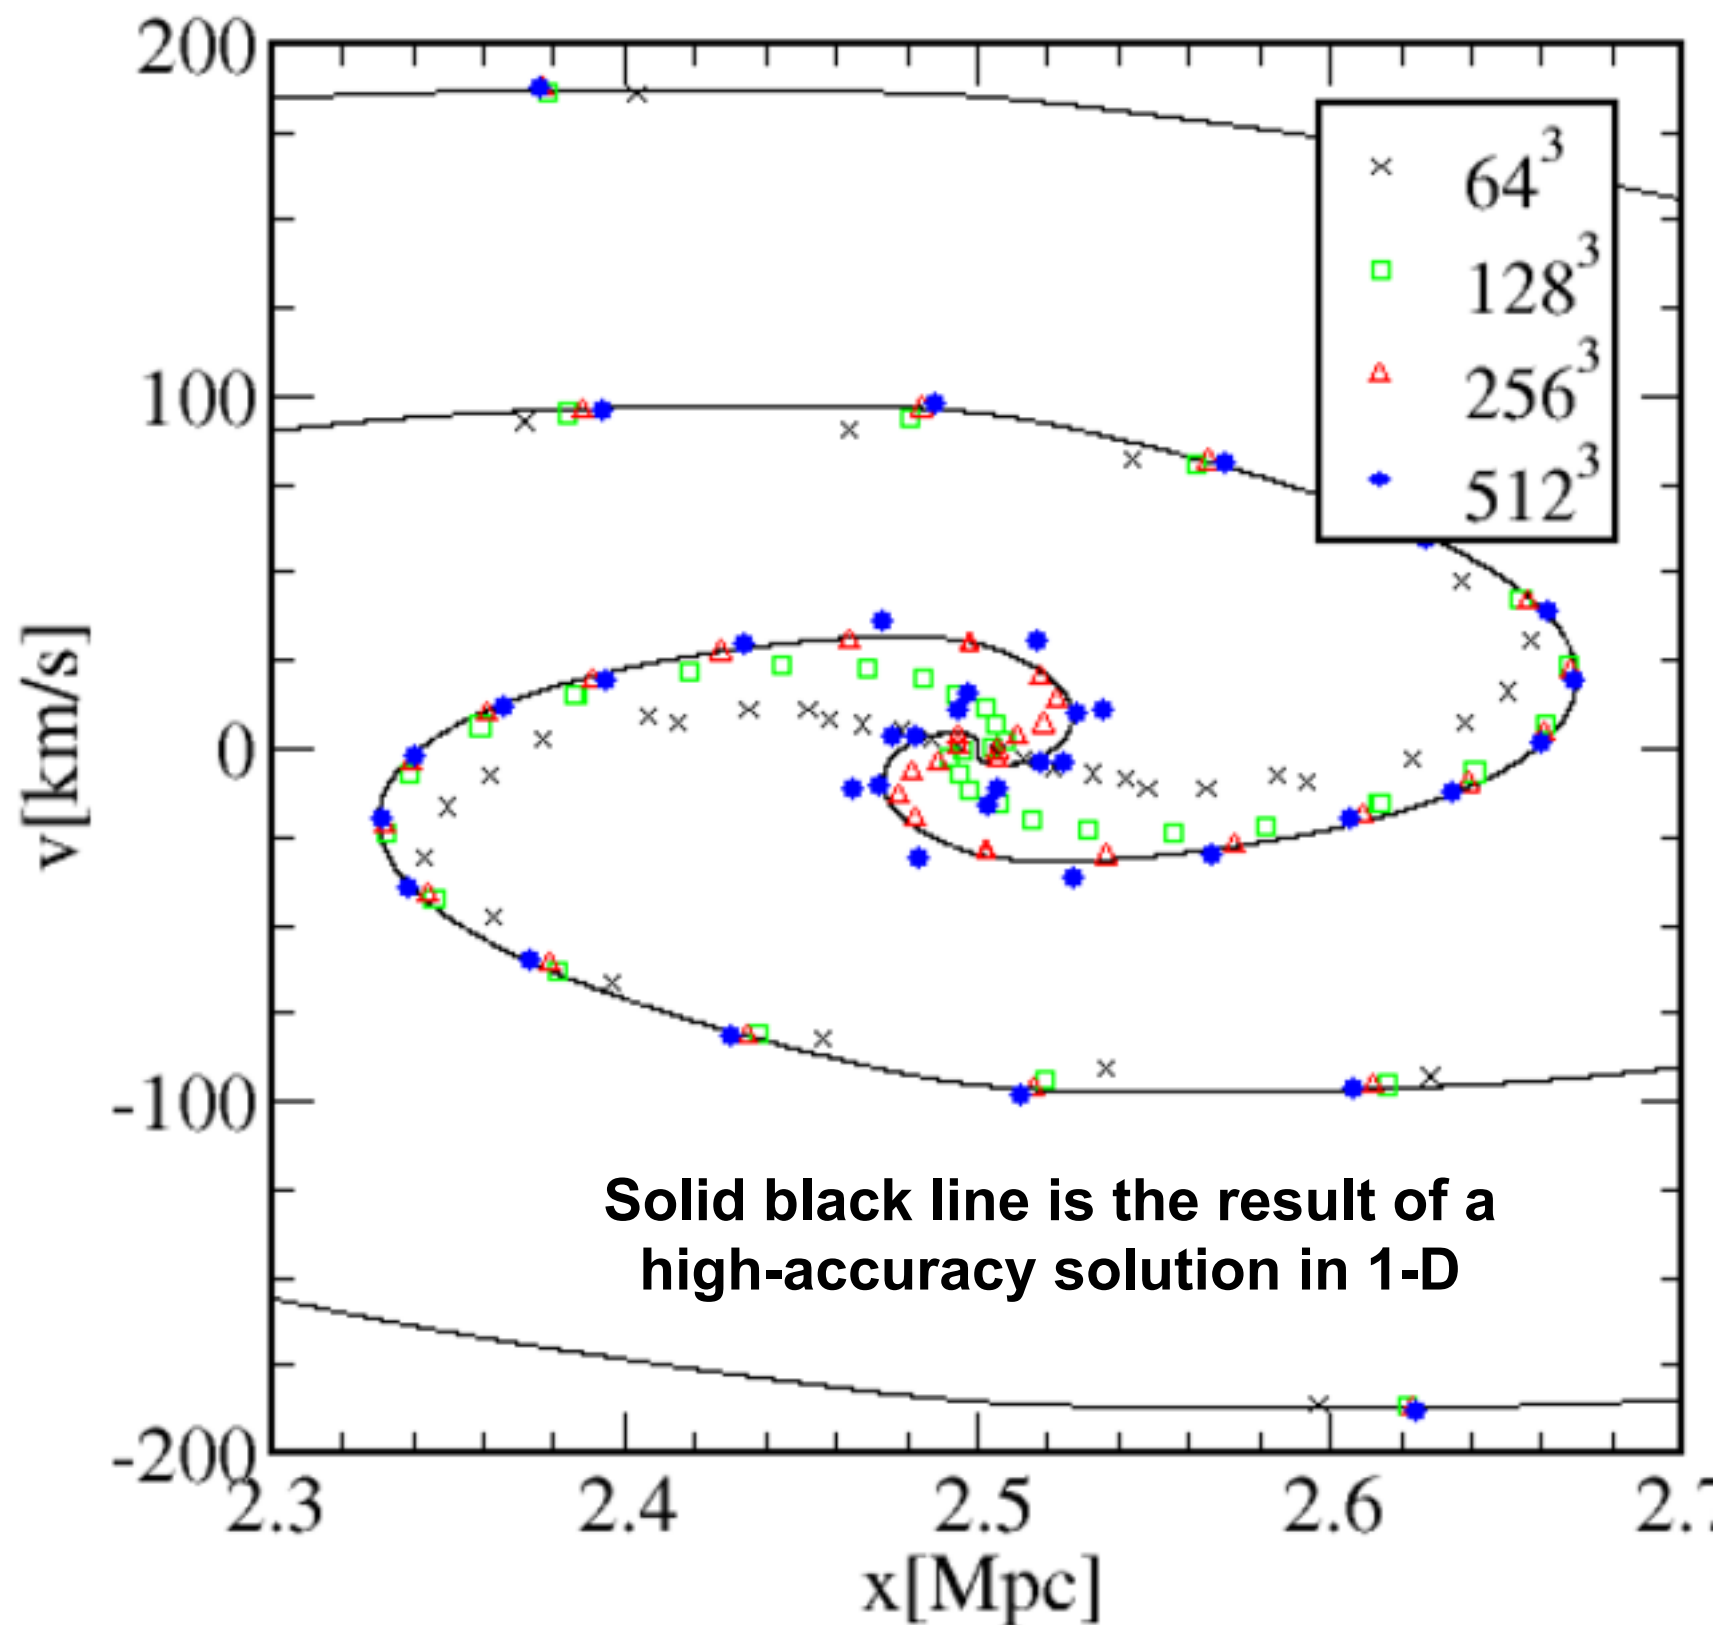
$$-\frac{L^2}{4\pi^2 k^2}$$

- This is the continuous Influence function
- Usually one obtains the potential and differentiates it with a 2nd order stencil and then interpolates the gradient on the particle with inverse CIC (preserves momentum)

# Case Study: PM (Planar Pancake Collapse Test)

- **Evidence of Collisionality**

- Run the Zeldovich pancake collapse test at multiple grid resolutions with particle number fixed
- Convergence must fail at some point (solution accuracy should improve for a while and then diverge)
- Convergence can be tracked until failure near the mid-plane at  $512^3$  grid points



# Case Study: PM (Higher-Order)

- **Higher-order Influence functions:**

- 4th: 
$$G_4(k) = \frac{3\Delta^2}{8} \frac{1}{\cos(2\pi k\Delta/L) - \frac{1}{16} \cos(4\pi k\Delta/L) - \frac{15}{16}}$$

- 6th:

$$G_6(k) = \frac{45\Delta^2}{128} \frac{1}{\cos(2\pi k\Delta/L) - \frac{5}{64} \cos(4\pi k\Delta/L) + \frac{1}{1024} \cos(8\pi k\Delta/L) - \frac{945}{1024}}$$

- These higher order functions should be used with appropriately smoothed density fields for their formal accuracy to be relevant
- The smoothing can be performed in Fourier space with a sinc-Gaussian filter (to isotropize the force)
- Higher-order gradients can be obtained directly in k-space using Super-Lanczos derivatives
- Time-stepping is done with symplectic integrators



# HACC PM Implementation

- **Spectral Particle-Mesh Solver:** Custom (large) FFT-based method -- uses 1) 6-th

$$G_6(\mathbf{k}) = \frac{45}{128} \Delta^2 \left[ \sum_i \cos \left( \frac{2\pi k_i \Delta}{L} \right) - \frac{5}{64} \sum_i \cos \left( \frac{4\pi k_i \Delta}{L} \right) + \frac{1}{1024} \sum_i \cos \left( \frac{8\pi k_i \Delta}{L} \right) - \frac{2835}{1024} \right]^{-1}$$



$$\left. \frac{\Delta f}{\Delta x} \right|_4 = \frac{4}{3} \sum_{j=-N+1}^N i C_j e^{(2\pi j x / L)} \frac{2\pi j \Delta}{L} \frac{\sin(2\pi j \Delta / L)}{2\pi j \Delta / L} - \frac{1}{6} \sum_{j=-N+1}^N i C_j e^{(2\pi j x / L)} \frac{2\pi j \Delta}{L} \frac{\sin(4\pi j \Delta / L)}{2\pi j \Delta / L}$$

where the  $C_j$  are the coefficients in the Fourier expansion of  $f$

$$S(k) = \exp \left( -\frac{1}{4} k^2 \sigma^2 \right) \left[ \left( \frac{2k}{\Delta} \right) \sin \left( \frac{k\Delta}{2} \right) \right]^{n_s}$$



$$f_{grid}(r) = \frac{1}{r^2} \tanh(br) - \frac{b}{r} \frac{1}{\cosh^2(br)} + cr (1 + dr^2) \exp(-dr^2) + e (1 + fr^2 + gr^4 + lr^6) \exp(-hr^2)$$

this later)



# Motivational Interlude: Hardware Evolution

- **Power is the main constraint**
  - Target: 30X performance gain by 2020
  - ~10-20MW per large system
  - Power/Socket roughly const.
- **Only way out: more cores**
  - Several design choices (e.g., cache vs. compute vs. interconnect)
  - All lead to more complexity
- **Micro-architecture gains sacrificed**
  - Accelerate specific tasks
  - Restrict memory access structure (SIMD/SIMT)
- **Machine balance sacrifice**
  - Memory/Flops; comm BW/Flops — all go in the “wrong” direction





# Emerging Architectures are Not New!

## ▶ HPC systems: “faster = more”

- More nodes
  - Separate memory spaces
  - Relatively slow network communication
- More complicated nodes
  - Architectures
    - Accelerators, multi-core, many-core
  - Memory hierarchies
    - CPU main memory
    - Accelerator main memory
    - High-bandwidth memory
    - Non-volatile memory

## ▶ Portable performance

- Massively parallel/concurrent
- Adapt to new architectures
  - Organize and deliver data to the right place in the memory hierarchy at the right time
  - Optimize floating point execution
- Not possible with off-the-shelf codes

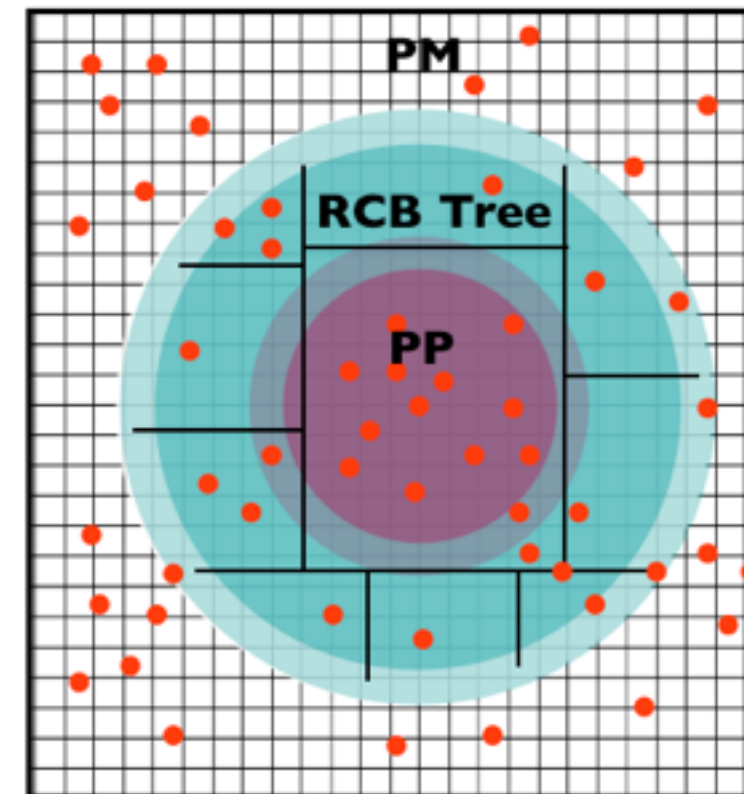


**Roadrunner Architecture (2008)**




# HACC: Design Principles

- **Optimize Code ‘Ecology’:** Numerical methods, algorithms, mixed precision, data locality, scalability, I/O, in situ analysis -- life-cycle significantly longer than architecture timescales
- **Framework design:** ‘Universal’ top layer + ‘plug-in’ optimized node-level components; minimize data structure complexity and data motion -- support multiple programming models
- **Absolute Performance:** Scalability, low memory overhead, and platform flexibility; minimal reliance on external libraries
- **Optimal Splitting of Gravitational Forces:** Spectral Particle-Mesh melded with direct and RCB tree force solvers, short hand-over scale (dynamic range splitting  $\sim 10,000 \times 100$ )
- **Compute to Communication balance:** Particle Overloading
- **Time-Stepping:** Symplectic, sub-cycled, locally adaptive
- **Force Kernel:** Highly optimized force kernel dominates compute time (90%), *no look-ups* due to short hand-over scale
- **Production Readiness:** runs on all supercomputer architectures; Gordon Bell Award Finalist 2012 and 2013, first production science code to break **10PFlops** sustained



**HACC force hierarchy  
(PPTreePM)**





# HACC Design: Portability Philosophy

- **Focus on Absolute Performance/Throughput:** High performance is a first-class requirement for HACC (portability here implies portability with *absolute* performance, not *relative* performance) — compute-intensive components control the performance of the code (not data motion)
- **Algorithmic Flexibility:** Allow for multiple algorithms in order to obtain the best possible performance on a given architecture (e.g., PPTreePM for BG/Q and Xeon Phi and P3M for GPUs)
- **Expert Tuning:** The code is designed so that the cost of obtaining performance is limited to experts tuning small subsets of node-level plug-in code (via particle overloading in HACC) — this is a *microkernel* based approach (the microkernel is specific to the code, not a general purpose routine)
- **Portable Top Layer:** Maximize portability of non-performance critical framework within which the compute-intensive kernels reside (in HACC, the spectral PM method is “soft” portable in this sense)
- **Limit External Dependencies:** Minimize reliance on non-vendor supported libraries that could impact performance, portability, and time to implementation on new platforms (the main HACC simulation path is entirely free of such libraries)




# Return from Interlude: 'HACC In Pictures'



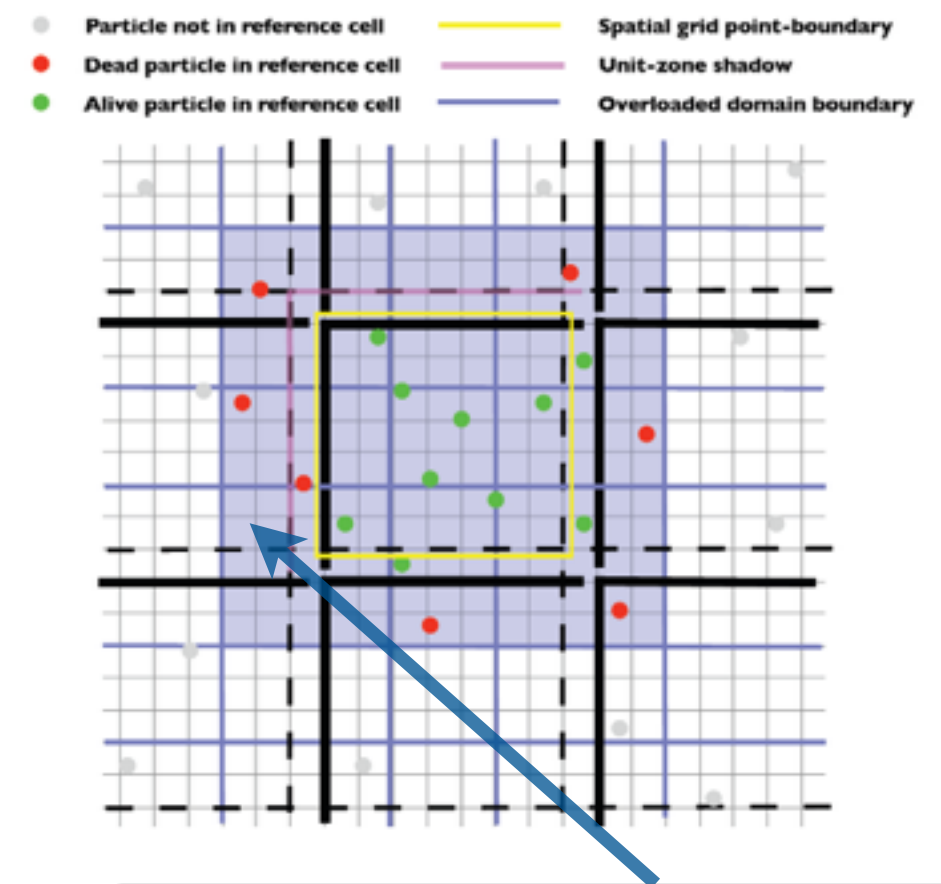


# Particle Overloading and Short-Range Solvers

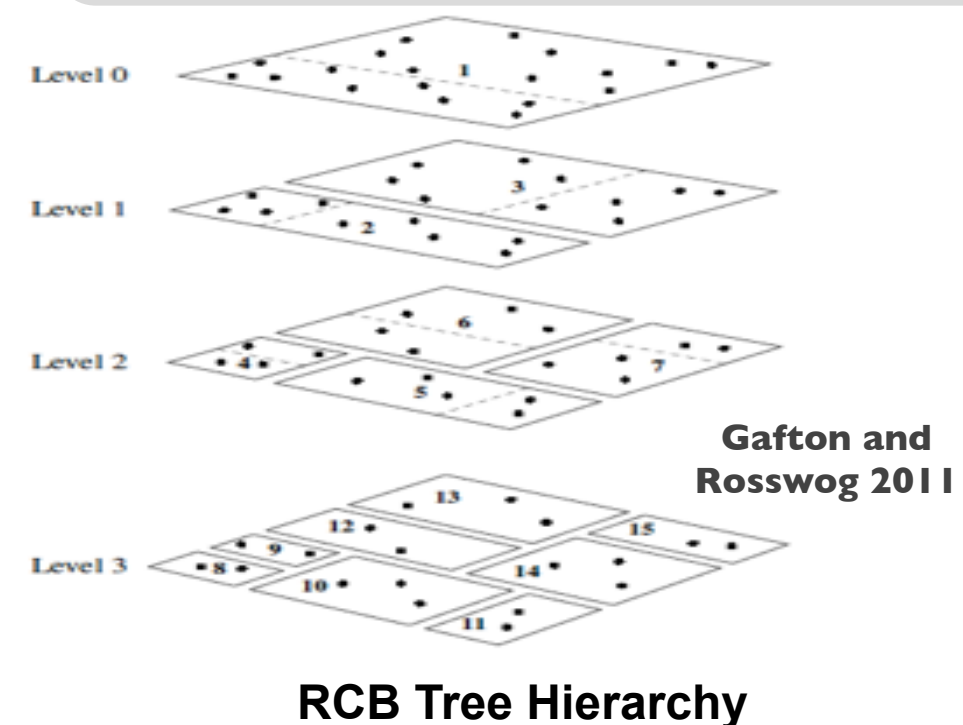
- **Particle Overloading:** Particle replication instead of conventional guard zones with 3-D domain decomposition -- minimizes inter-processor communication and allows for swappable short-range solvers (**IMPORTANT**)
- **Short-range Force:** Depending on node architecture switch between P3M and PPTreePM algorithms (pseudo-particle method goes beyond monopole order), by tuning number of particles in leaf nodes and error control criteria, optimize for computational efficiency
- **Error tests:** Can directly compare different short-range solver algorithms
- **Load-balancing:** Passive + Active task-based



HACC Force Algorithm Test: PPTreePM vs. P3M



Overload Zone (particle 'cache')

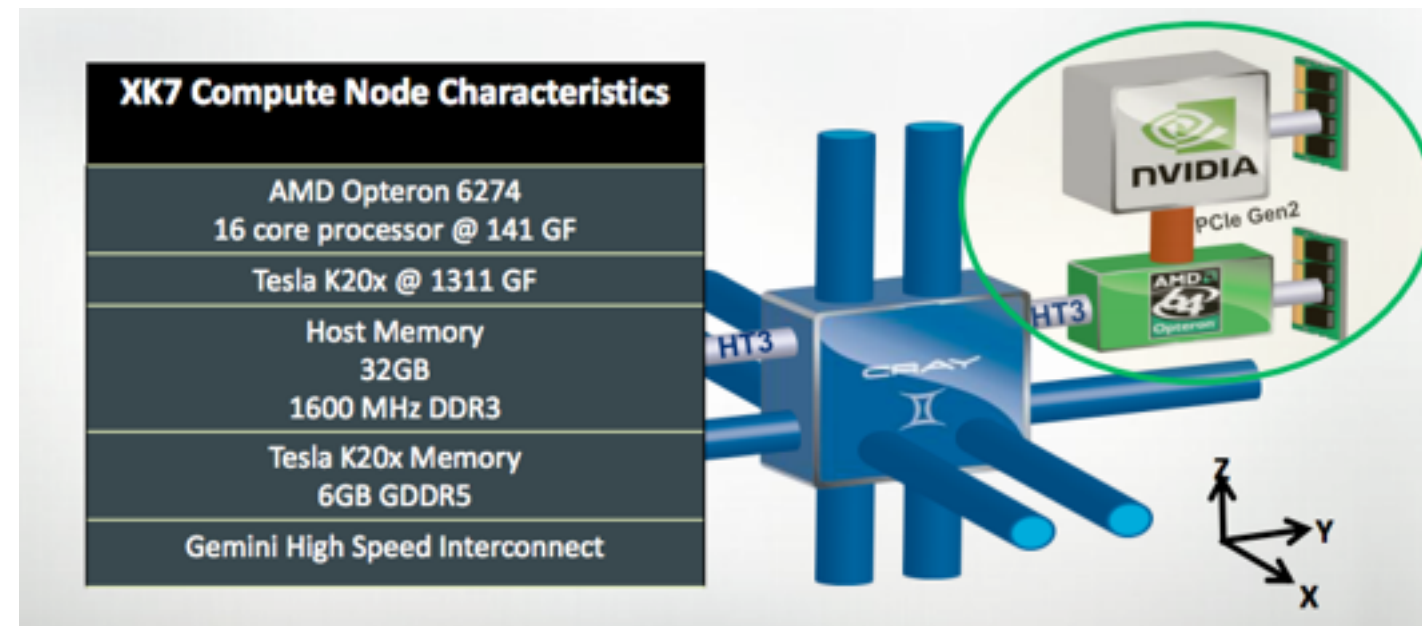


RCB Tree Hierarchy

# Accelerated Systems: Specific Issues

## Imbalances and Bottlenecks


- Memory is primarily host-side (32 GB vs. 6 GB) (against Roadrunner's 16 GB vs. 16 GB), important thing to think about (in case of HACC, the grid/particle balance)
- PCIe is a key bottleneck; overall interconnect B/W does not match Flops (not even close)
- There's no point in 'sharing' work between the CPU and the GPU, performance gains will be minimal -- GPU must dominate
- The only reason to write a code for such a system is if you can truly exploit its power (2 X CPU is a waste of effort!)



## Strategies for Success

- It's (still) all about understanding and controlling data motion
- Rethink your code and even approach to the problem
- Isolate hotspots, and design for portability around them (modular programming)
- Like it or not, pragmas will never be the full answer

# HACC on Titan: GPU Implementation (Schematic)



## P3M Implementation (OpenCL & CUDA)      New Implementations/Improvements

- 1D-decomposed data pushed to GPU in large blocks; data sub-partitioned into chaining-mesh cubes
  - Compute inter-particle forces within cubes and neighboring cubes
  - Large block size ensures computational time far exceeds memory transfer latency
  - Natural parallelism provides high performance wrt book-keeping required for tree algorithms
- P3M data-push once every long time-step, with 'soft boundary' chaining mesh, completely eliminates latency
  - TreePM analog of BG/Q code written in CUDA also provides high performance
  - Each block is an independent work-item; timing the blocks is used in a load-balancing scheme — blocks are transferred to lightly loaded ranks during execution

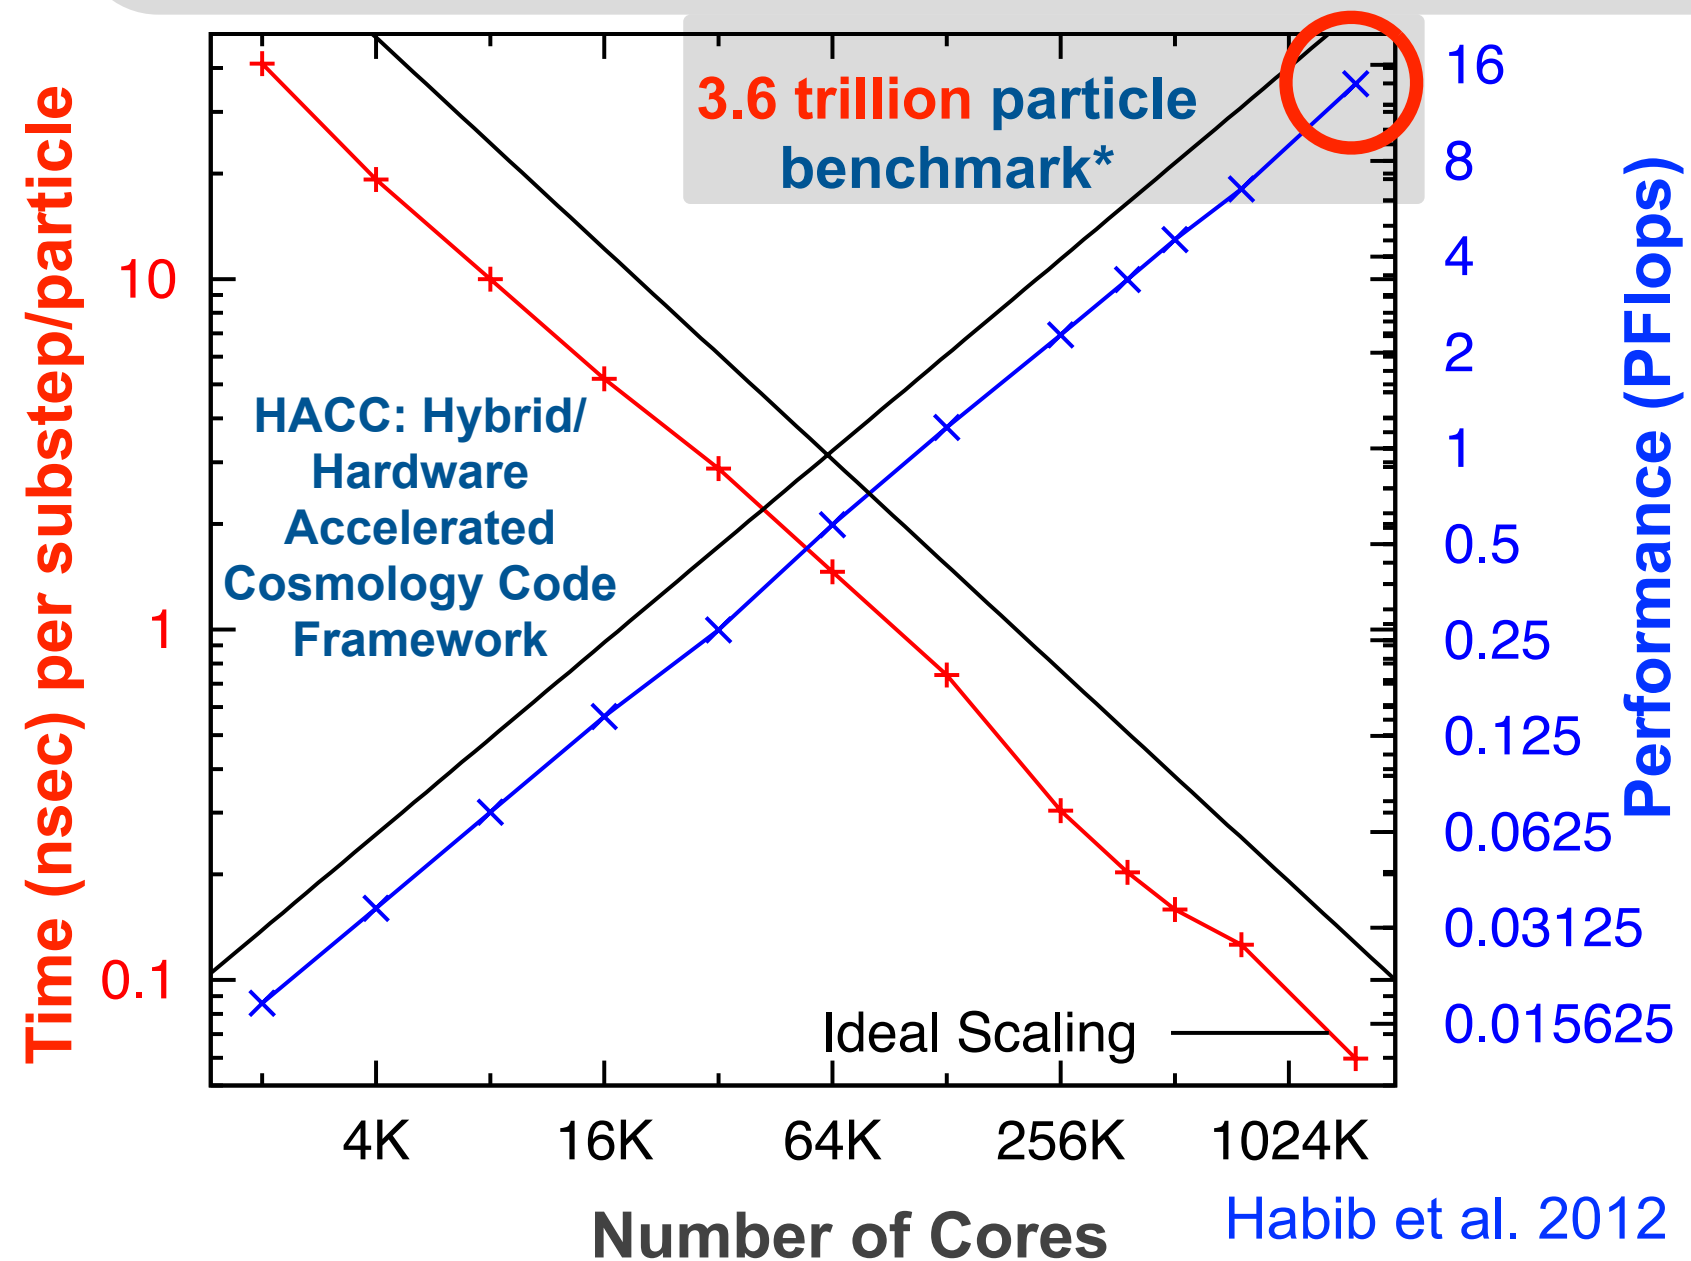


# HACC on the BG/Q ('pre-manycore')

## HACC BG/Q Version

- **Algorithms:** FFT-based SPM; PP+RCB Tree
- **Data Locality:** Rank level via 'overloading', at tree-level use the RCB grouping to organize particle memory buffers
- **Build/Walk Minimization:** Reduce tree depth using rank-local trees, shortest hand-over scale, bigger p-p component
- **Force Kernel:** Use polynomial representation (no look-ups); vectorize kernel evaluation; hide instruction latency

**13.94 PFlops, 69.2% peak, 90% parallel efficiency on 1,572,864 cores/MPI ranks, 6.3M-way concurrency**




**\*largest ever run**

**HACC weak scaling on the IBM BG/Q (MPI/OpenMP)**


# HACC on Titan: GPU Implementation Performance

- P3M kernel runs at 1.6TFlops/node at 40.3% of peak (73% of algorithmic peak)
- TreePM kernel was run on 77% of Titan at 20.54 PFlops at almost identical performance on the card
- Because of less overhead, P3M code is (currently) faster by factor of two in time to solution




99.2% Parallel Efficiency

**MC<sup>3</sup> (HACC precursor)**



**Gadget-2**



**Snapshot from Code Comparison simulation, ~25 Mpc region; halos with > 200 particles,  $b=0.15$**   
**Differences in runs: P<sup>3</sup>M vs. TPM, force kernels, time stepper: MC<sup>3</sup>: a; Gadget-2:  $\log(a)$**   
**Power spectra agree at sub-percent level**

